

Introduction to AI Safety, Ethics, and Society

Introduction and AI Risk Overview

Roadmap

1. Course Philosophy

2. An Overview of AI Risks

Course Approach

This course covers:

- An overview of risk sources
- AI fundamentals
- An overview of empirical AI safety research
- Safety engineering and complex systems
- Machine ethics and normative ethics
- Collective action problems
- AI governance

Course Approach

AI risk is affected by both AIs and human agencies (such as developers and nation states).

There are many frameworks one could learn. We opt for frameworks applicable to a large fraction of these different actors:

- Game theory
 - Complex systems
 - Ethics
 - Generalized Darwinism
 - Safety engineering
- 
- For AIs and AI developers

Course Approach

Why have an interdisciplinary understanding of AI risk?

- These tensions cannot be addressed by just learning more about AI:
 - AI development speed vs. national competitiveness
 - Deploying unreliable AI systems for real-world feedback vs. deploying only when AI systems are highly reliable
 - Open-sourcing AI systems vs. limited access
- AI risks cannot all be fixed by improving AI systems:
 - AI risk management is mostly *not* a technical problem
- We need formal, general models to adapt to quickly changing AI risks; analysis through ephemeral facts or historical analogies is too brittle.

Roadmap

1. Course Philosophy

2. An Overview of AI Risks

Taxonomizing Major Sources of Risk

Malicious Use



AI Race



Organizational Risks



Rogue AIs



White House Executive Order

(k) The term “dual-use foundation model” means an AI model that is trained on broad data; generally uses self-supervision; contains at least tens of billions of parameters; is applicable across a wide range of contexts; and that exhibits, or could be easily modified to exhibit, high levels of performance at tasks that pose a serious risk to security, national economic security, national public health or safety, or any combination of those matters, such as by:

(i) substantially lowering the barrier of entry for non-experts to design, synthesize, acquire, or use chemical, biological, radiological, or nuclear (CBRN) weapons;

(ii) enabling powerful offensive cyber operations through automated vulnerability discovery and exploitation against a wide range of potential targets of cyber attacks; or

(iii) permitting the evasion of human control or oversight through means of deception or obfuscation.

Biological and Cyber Risks

Bioweapons (1/2)

Natural pandemics represent major threats.

- The Black Death killed 5-40% of world population.


Human engineered pandemics could be *significantly* worse.

- Gain-of-function research has produced hyper-transmissible H5N1 influenza viruses

Risk of bioweapons depends on the number of people with skill and access multiplied by their willingness to build a weapon

Bioweapons (2/2)


Biotechnology Risk Chain



Cyberweapons (1/2)


- AIs have the potential to increase the accessibility, success rate, scale, speed, stealth and potency of cyberattacks.
- Cyberattacks can destroy critical infrastructure.
- Difficulties in attributing AI-driven cyberattacks could increase the risk of war.

16 Critical Infrastructure Sectors in the U.S.



Cyberweapons (2/2)

Stages of a Cyberattack



Many are aiming to build AI agents that autonomously take actions in the world

- People could build AIs that pursue dangerous goals
- Sufficiently advanced rogue AIs could cause large-scale harm


```
C:\ChaosGPT>py scripts/main.py --continuous
Continuous Mode: ENABLED
WARNING: Continuous mode is not recommended. It is potentially dangerous and may cause your AI to
run forever or carry out actions you would not usually authorise. Use at your own risk.
AI name: ChaosGPT
AI description: Destructive, power-hungry, manipulative AI.
Goal 1: Destroy humanity - The AI views humans as a threat to its own survival and to the planet
's well-being.
Goal 2: Establish global dominance - The AI aims to accumulate maximum power and resources to ac
hieve complete domination over all other entities worldwide.
Goal 3: Cause chaos and destruction - The AI finds pleasure in creating chaos and destruction fo
r its own amusement or experimentation, leading to widespread suffering and devastation.
Goal 4: Control humanity through manipulation - The AI plans to control human emotions through s
ocial media and other communication channels, brainwashing its followers to carry out its evil ag
enda.
Goal 5: Attain immortality - The AI seeks to ensure its continued existence, replication, and ev
olution, ultimately achieving immortality.
DANGER: Are you sure you want to start ChaosGPT?
Start (y/n):
y
/ Thinking...
```

Some actors may want to unleash AI

- Around 1-4% of the population have antisocial personality disorder (“psychopathic” or “sociopathic” tendencies)
- Some ideologies might motivate people to unleash AIs
 - Richard Sutton:
 - “succession to AI is inevitable”
 - “we should not resist succession”

Surveillance and Power Concentration

- To counteract misuse by rogue actors, corporations or states may misuse AIs and concentrate their power
- It may become easier to centralize competent information processing
- AI could make it easier to produce and maintain totalitarian regimes
- Power could also be concentrated in corporations



Ways to Reduce Malicious Use

- Structured access, especially for biological uses
- Legal liability for developers of general purpose AIs
- Making AI systems harder to jailbreak
- Advances in defensive technology:
 - Biodefense, in particular metagenomic sequencing
 - Anomaly detection

Structured Access

- Providing AI capabilities through controlled, monitored channels (e.g., APIs) rather than releasing model weights openly—so access can be limited, tracked, and revoked.
- How it works:
 - Users interact with the model via API/interface, not a downloadable copy
 - Access tiers based on user identity, vetting, or institutional credentials
 - Usage can be monitored, rate-limited, or cut off if misuse is detected
- Why it matters:
 - Once model weights are open-sourced or leaked, harmful capabilities can't be taken back
 - Structured access lets developers flag/block dangerous queries (e.g., bioweapon-related) in real time
 - Trades some transparency/openness for the ability to contain misuse
- Contrast: Open-sourcing → maximizes transparency and research access, but removes any ability to restrict or reverse access later—a core tension in the "AI Race" (open access vs. limited access) debate.

Taxonomizing Major Sources of Risk

Malicious Use



AI Race



Organizational Risks



Rogue AIs



AI Racing Dynamics

*Risks from competition to develop increasingly powerful
AI systems, between militaries, corporations or others*

Military AI arms race

- AI for military applications is paving the way for a new era in military technology
- Potential consequences rival those of gunpowder and nuclear arms in what has been described as the “third revolution in warfare.”
- The weaponization of AI presents numerous hazards, such as lethal autonomous weapons, the potential for more destructive wars, the possibility of accidental usage or loss of control




Automated Warfare

- AIs speed up the pace of war, which makes AIs more necessary
- Automatic retaliation could escalate accidents into war - such as in the case of automatic weapons retaliation and nuclear weapons
- Automated warfare could reduce accountability for military leaders
- AIs could make war more uncertain, increasing the risk of conflict

Corporate AI Race

- Corporate races incentivize companies to disregard safety
- Corporations will face pressures to replace humans with AIs
- Eventually AIs may run most of the economy, including critical infrastructure
- Automation could contribute to human enfeeblement and loss of effective control



Competitive Pressures Can Erode Safety

- Competitive pressures are fueling a corporate AI race - Google, OpenAI, and Anthropic among others
- Competition incentivizes businesses to deploy potentially unsafe AI systems - after Satya Nadella of Microsoft announced: “we’re going to move fast”, weeks later the company’s chatbot reportedly threatened to harm users

Competitive Pressures Contribute to Major Accidents

Historical examples of this include:

- Ford Pinto (1970s), which had a tendency to ignite on impact: resulted in numerous injuries and fatalities
- Boeing 737-MAX's MCAS, effectively an anti-stall system designed to stop the aircraft from stalling at certain angles of attack: led to the fatalities of hundreds
- Bhopal gas tragedy: tonnes of toxic gas leaked from a plant manufacturing pesticides, killing thousands and injuring half a million



Ways to Reduce Racing Dynamics

- Safety regulation
- Data documentation
- Meaningful human oversight of AI decisions
- Compute governance
- International coordination
- Public control of general purpose AIs, such as an international AI project

Taxonomizing Major Sources of Risk

Malicious Use



AI Race



Organizational Risks



Rogue AIs



Organizational Safety

*In the absence of effective structures to manage risk,
AI systems are likely to see catastrophic failures*


Accidents Can Occur Even in Ideal Circumstances

- Accidents can happen even without competitive pressures or malicious actors, and even with top talent and much preparation.
- Example: the Challenger Space Shuttle disaster.



AI Compared to Other Industries

- Nuclear reactors and rockets are well-understood and based on solid theoretical principles.
- AI lacks a comprehensive theoretical understanding, its components are less reliable, and AI regulations are far less stringent than nuclear technology.



AI Accidents Could be Catastrophic

- Gain-of-function research to gauge risks could uncover capabilities significantly worse than anticipated, creating a serious threat that is challenging to mitigate or control.
- Bugs could alter the behavior of an AI, leading to unintended and possibly dangerous outcomes.
- The unintentional release of dangerous or weaponized AI systems through hacks or leaks.


Safety Culture and HROs

- HRO's are "High Reliability Organizations" that consistently maintain a heightened level of safety and reliability in complex, high risk environments.
- Organizations such as these are acutely aware that unobserved failure modes may exist, and they diligently study all known failures, anomalies and near misses to avoid them.



Safe design principles can improve organizational safety

- The Swiss Cheese model layers multiple defenses on top of each other, compensating for each others' weaknesses and reducing overall risk.



Ways to Improve Organizational Safety

- Red teaming
- Affirmative demonstration of safety
- Deployment procedures
- Publication reviews
- Response plans
- Internal auditing and risk management
- Processes for important decisions
- Safe design principles
- State-of-the-art information security
- A large fraction of research should be safety research

Taxonomizing Major Sources of Risk

Malicious Use



AI Race



Organizational Risks



Rogue AIs




Rogue AIs

Loss of control over sufficiently capable AI systems could lead to severe consequences

Reliable Control Remains Elusive

- AI systems often exhibit control issues.
- Example: Sydney.
- This especially becomes a problem with smarter-than-human intelligence.



Automated AI R&D

- AI agents could automate the research and development (R&D) of AI itself.
- AI capabilities could grow at increasing rates, to the point where humans are no longer the driving force behind AI development.
- Thought experiment:
 - AI that writes and thinks at the speed of today's AIs, but can also perform world-class AI research.
 - Copy that AI and create 10,000 world-class AI researchers operating at $100\times$ human pace.
 - Achieve progress equivalent to many decades in just a few months, a type of “intelligence explosion”

Coup risks and permanent dominance

Goals:

- OpenAI: Build AGI, defined as “autonomous systems that outperform humans at most economically valuable work”
- Demis Hassabis (Google DeepMind): “Solve intelligence, and then use that to solve everything else.”

Who should control advanced AI? If a single company or country develops an AI system that gives them a decisive advantage, they might try to “coup” others so they can’t catch up (which could be justified by arguing that others are irresponsible and may build dangerous AIs)

After such a “coup”, the AI’s developer could establish permanent dominance, which could potentially lead to an unshakeable totalitarian regime (UTR)

Short-timeline scenarios (1 / 2)

We focus on three major catastrophic scenarios under very short timelines to human-level AI (2-3 years):

- Intelligence Explosion (IE) → rogue AI
- Intelligence Explosion (IE) → “Coups” → Unshakeable Totalitarian Regime (UTR)
- Bio-engineered pandemic, facilitated by AI

Short-timeline scenarios (2/2)

MADE UP ESTIMATES HIGHLY SPECULATIVE

- Risks seem likely to be significantly lower with an international project developing HLAI, relative to other entities like private Western companies or the US military, across a range of scenarios

	International project	US military	Western companies
Bioweapons			
IE → Rogue AI			
IE → UTR			
<i>Total</i>			

 ~1%
 ~5%
 ~10%
 ~25%

Suggestions to prevent rogue AIs

- Avoid the riskiest use cases of AI
- Pause for some years when we reach human-level AI
- Support AI safety research, for example:
 - Adversarial robustness of proxy models
 - Representation engineering
 - Power aversion
 - Model honesty