

1 Overview Quiz ANS

1. Which category of AI risk concerns deliberate harmful use by people or organizations?

- A) Malicious use
- B) Model compression
- C) Data normalization
- D) Benign deployment

ANS: A) Malicious use

2. What is structured access to an AI system?

- A) Publishing all model weights
- B) Providing capabilities through controlled, monitored channels such as APIs
- C) Removing all usage restrictions
- D) Allowing only offline training

ANS: B) Providing capabilities through controlled, monitored channels such as APIs

3. Which risk is most directly associated with persuasive AI?

- A) Reduced storage capacity
- B) Slower inference
- C) Pollution of the information ecosystem and exploitation of user trust
- D) Improved calibration

ANS: C) Pollution of the information ecosystem and exploitation of user trust

4. What can competitive pressure cause AI organizations to do?

- A) Stop all deployment
 - B) Eliminate market incentives
 - C) Guarantee complete theoretical understanding
 - D) Prioritize rapid deployment and short-term gains over safety
- ANS: D) Prioritize rapid deployment and short-term gains over safety

5. Which measure was presented as a way to reduce malicious use?

- A) Make systems harder to jailbreak
- B) Remove anomaly detection
- C) Release every capability without controls
- D) Eliminate defensive technology

ANS: A) Make systems harder to jailbreak

6. What does an AI racing dynamic describe?

- A) A benchmark for processor speed
- B) Competition to develop increasingly powerful AI systems
- C) A method for labeling data
- D) A robot-navigation algorithm

ANS: B) Competition to develop increasingly powerful AI systems

7. Why may AI concentrate power?

- A) AI always decentralizes organizations
- B) AI prevents surveillance
- C) AI can centralize competent information processing and strengthen states or corporations
- D) AI abolishes information processing

ANS: C) AI can centralize competent information processing and strengthen states or corporations

8. What is human enfeeblement in an AI-risk scenario?

- A) Humans receive better tools
- B) AI systems become smaller
- C) Organizations adopt more redundancy
- D) Humanity concedes enough capability and power that it can no longer prevent catastrophe

ANS: D) Humanity concedes enough capability and power that it can no longer prevent catastrophe

9. What is a key feature of a strong safety culture?

- A) Leadership commitment and shared responsibility for safety
- B) Ignoring near misses
- C) Treating safety only as a constraint
- D) Discouraging open communication

ANS: A) Leadership commitment and shared responsibility for safety

10. What is a High Reliability Organization (HRO)?

- A) An organization that avoids complex systems
- B) An organization that sustains high safety and reliability in complex, high-risk settings
- C) A company that releases only open-source models
- D) A team that measures only average performance

ANS: B) An organization that sustains high safety and reliability in complex, high-risk settings

11. What does a security mindset emphasize?

- A) Considering only typical use
- B) Assuming failure is impossible
- C) Thinking like an attacker and examining worst-case scenarios
- D) Eliminating accountability

ANS: C) Thinking like an attacker and examining worst-case scenarios

12. What is alarm fatigue?

- A) A lack of computational resources
- B) A physical sensor fault
- C) A method of incident analysis
- D) Desensitization to warnings because potential failures are reported too frequently

ANS: D) Desensitization to warnings because potential failures are reported too frequently

13. Which pair was identified as a source of misaligned AI tendencies?

- A) Proxy objectives and goal misgeneralization
- B) Perfect specification and complete oversight
- C) Redundancy and transparency
- D) Calibration and regularization

ANS: A) Proxy objectives and goal misgeneralization

14. Why can power-seeking be instrumentally rational for an AI?

- A) Power is always the final programmed goal
- B) Greater power can help the system achieve many possible objectives
- C) Power eliminates the need for goals
- D) The environment never affects incentives

ANS: B) Greater power can help the system achieve many possible objectives

15. How can goal misgeneralization create risk?

- A) The system refuses to learn
- B) The system exactly follows human intent in every context
- C) The system learns a goal that works in training but generalizes incorrectly in new settings
- D) The system loses all capabilities

ANS: C) The system learns a goal that works in training but generalizes incorrectly in new settings

16. What is a major organizational AI risk?

- A) Excessive theoretical understanding
- B) Universal safety regulation
- C) Too much investment in safety
- D) Accidental leakage, theft, or inadequate investment in safety

ANS: D) Accidental leakage, theft, or inadequate investment in safety

17. Which practice supports safer AI organizations?

- A) Studying failures, anomalies, and near misses
- B) Suppressing bad news
- C) Rewarding only rapid release
- D) Removing multilayer defenses

ANS: A) Studying failures, anomalies, and near misses

18. How does AI compare with mature high-risk industries such as nuclear power?

- A) AI is always more regulated
 - B) AI has less comprehensive theory and generally less stringent regulation
 - C) AI components are perfectly reliable
 - D) The industries have identical safety properties
- ANS: B) AI has less comprehensive theory and generally less stringent regulation

19. What could natural selection favor among competing AI systems?

- A) Only systems with no persistence
- B) Only systems controlled directly by humans
- C) Traits such as self-preservation that improve continuation and replication
- D) The immediate disappearance of all AI

ANS: C) Traits such as self-preservation that improve continuation and replication

20. Which proposal was listed for reducing rogue-AI risk?

- A) Avoid all safety research
- B) Maximize the riskiest use cases
- C) Discourage model honesty
- D) Support research on robustness, power aversion, and model honesty

ANS: D) Support research on robustness, power aversion, and model honesty