

6 Machine Ethics Quiz

1. Which description fits "standards" as the deck contrasts them with "rules" in law?

- A) Laws written specifically to govern AI systems only
- B) Flexible statements that allow proportional responses
- C) Informal moral intuitions that carry no legal force at all
- D) Rigid instructions that reduce interpretation problems

ANS:

2. What is "gaming" in the discussion of law and AI?

- A) AIs taking short, open-ended instructions too literally and misreading what was meant
- B) AIs refusing to follow any of the rules that they have been given by developers
- C) Developers testing a model against a standard benchmark before its public release
- D) AIs playing the system of rules to their advantage in undesirable ways

ANS:

3. Which two problems does the deck give for relying on law alone to guide AI?

- A) Law is too moral in its content, and it ignores economic concerns
- B) Law is too specific to apply, and it changes too quickly to enforce
- C) Law is kept secret from the public, and it applies only to corporations
- D) Law is silent in many areas, and some laws clash with human morals

ANS:

4. How does the deck define "reasonable care"?

- A) The caution an ordinary, prudent person would use to avoid foreseeable harm to others
- B) The maximum level of caution that is technically possible for a company to achieve
- C) Following every instruction that a user gives to the system without any exception
- D) Avoiding every possible source of legal liability that the company could face

ANS:

5. What was the central dispute about the COMPAS software?

- A) Alleged excessive accuracy versus the counterclaim that the software was too cheap to maintain
- B) Alleged use only in healthcare versus the counterclaim that it was used only in schools
- C) Alleged bias against African Americans versus the counterclaim that it was calibrated and therefore fair
- D) Alleged copyright violation versus the counterclaim that the software was open source

ANS:

6. Which fairness definition requires the false positive rate and false negative rate to be equal across groups?

- A) Counterfactual fairness
- B) Equalized odds
- C) Statistical parity
- D) Calibration

ANS:

7. What does the impossibility theorem for AI fairness state?

- A) No AI system can ever be allowed to make a decision that affects a person
- B) Fairness can always be achieved simply by collecting a great deal more training data
- C) The three classic definitions of fairness are usually contradictory, so no predictor satisfies them all
- D) All of the classic fairness definitions turn out to be logically equivalent to each other

ANS:

8. According to the deck, which is a weakness of social approaches to improving fairness (such as participatory design)?

- A) They cannot involve the stakeholders affected
- B) They fail to address any bad outcomes at all
- C) They are more subjective, expensive, and difficult
- D) They work only for a single narrow domain

ANS:

9. Which is among the conditions listed for the First Fundamental Theorem of Welfare Economics?

- A) A single seller is big enough to set prices alone
- B) Externalities are common in consumption
- C) The state does not enforce contracts
- D) Everyone has perfect information

ANS:

10. What does "information asymmetry" describe, and what example does the deck give?

- A) Both parties have identical information, as in a perfect market
- B) A firm with enough influence to set the market price, as with Standard Oil
- C) One party exploits knowledge the other lacks, as in predatory lending
- D) Consequences that affect unrelated third parties, as with pollution

ANS:

11. Into how many companies did the US break up Standard Oil, the deck's example of addressing monopolies?

- A) 12
- B) 7
- C) 34
- D) 50

ANS:

12. Which of the following is listed as a way to address or internalize externalities?

- A) Litigation and liability, defined property rights, and taxation
- B) Mergers, rationing, and censorship of market information
- C) Secret contracts, price fixing, and the formation of monopolies
- D) Open weights, subsidies, and aggressive advertising campaigns

ANS:

13. What does the Gini coefficient measure?

- A) A nation's total economic output over one year
- B) The number of firms competing within a market
- C) The average life expectancy of the population
- D) How a nation's income or wealth is distributed

ANS:

14. Why does the deck say that GDP is not social wellbeing?

- A) It counts only those goods and services that people say they dislike
- B) It always falls whenever the level of automation in the economy rises
- C) It directly measures how happy and satisfied the whole population is
- D) Many socially valuable tasks, like parenting, are not captured by it

ANS:

15. Which account of wellbeing is also called hedonism?

- A) Wellbeing as a set of objective goods such as knowledge
- B) Wellbeing as the balance of pleasure over pain
- C) Wellbeing as the satisfaction of preferences and desires
- D) Wellbeing as the rate of economic growth

ANS:

16. Which kind of preference is simply what people choose, and what is one of its drawbacks?

- A) Revealed preference; weakness of will
- B) Stated preference; weakness of will
- C) Idealized preference; manipulation by others cannot occur
- D) Revealed preference; it always reflects fully informed choices

ANS:

17. What are idealized preferences?

- A) What a fully informed person would choose
- B) What people say their preferences are
- C) What people actually choose in practice
- D) What an AI system predicts the average person prefers

ANS:

18. What is the "wireheading" problem in the happiness section?

- A) Happiness cannot be represented by any kind of mathematical function
- B) An AI might stimulate people into happiness without regard to what they value
- C) AIs refuse to take any action that could improve the happiness of humans
- D) People become happier only when their income increases in a given year

ANS:

19. In the discussion of moral uncertainty, what does "credence" mean?

- A) The legal authority that a moral theory holds in the courts of a country
- B) The probability an individual assigns to a moral theory being true
- C) The amount of utility that an action produces for the people it affects
- D) The number of people in a society who currently endorse a particular theory

ANS:

20. What does the Moral Parliament approach do?

- A) Treating the decision as a negotiation among several moral views
- B) Supporting only the theory in which one has the highest credence
- C) Letting citizens vote on every individual AI decision
- D) Computing the highest average moral value under one single theory

ANS: