

8 Governance Quiz ANS

1. Which of the following is an example of a financial (dis)incentive among the governance tools in the deck?

- A) Maintaining an AI chip registry for the government
- B) Taxes, liability, and export controls
- C) A summit that raises awareness among countries
- D) NIST and ISO standards for good practice

ANS: B) Taxes, liability, and export controls

2. What mechanism does the deck propose for AI-driven explosive economic growth?

- A) AIs substituting for human labor lifts a productivity bottleneck, creating a positive feedback loop
- B) Cutting the number of workers in the whole economy so that wages for those left rise
- C) Banning automation outright so that existing human jobs are preserved and growth stays slow
- D) Humans substituting for machines raises productivity, creating a negative feedback loop

ANS: A) AIs substituting for human labor lifts a productivity bottleneck, creating a positive feedback loop

3. What is Baumol's cost disease in the discussion of bottlenecks on AI automation?

- A) Automation always lowers the cost of every kind of job across the whole economy
- B) Non-automated jobs may see rising costs, offsetting the overall economic impact
- C) Wages fall to zero in every sector of the economy once AI is introduced widely
- D) Energy becomes unlimited, so every bottleneck on AI automation disappears altogether

ANS: B) Non-automated jobs may see rising costs, offsetting the overall economic impact

4. Why does the deck say that human-level AI could create an unequal distribution of benefits?

- A) Those who own supercomputers capture the benefits of automating labor
- B) Everyone receives the same share of the new wages that automation produces
- C) Workers are paid more and more as their tasks become automated by AI systems
- D) Computing power is distributed equally to every household, so benefits are shared

ANS: A) Those who own supercomputers capture the benefits of automating labor

5. Which approach does the deck describe as providing structured access to AI?

- A) Releasing the model weights to the whole public openly, with no restrictions at all
- B) Keeping the model private, strictly within the one company that trained it
- C) Offering the model through APIs with restrictions, possibly with KYC
- D) Deleting the model weights immediately after training so that nobody can use them

ANS: C) Offering the model through APIs with restrictions, possibly with KYC

6. Why does the deck describe releasing model weights as irreversible proliferation?

- A) Released weights cannot be taken back, and their uses cannot be controlled
- B) Weights stop working a short time after a release, so the risk expires quickly
- C) Releasing weights always reduces risk in the long run, because many people can audit them
- D) Only governments are able to run the released weights on their own hardware

ANS: A) Released weights cannot be taken back, and their uses cannot be controlled

7. Which risk does the deck associate with tightly controlled models?

- A) Bottom-up malicious use by rogue individuals and nonstate actors with no authority
- B) Excessive transparency that exposes the model's inner workings to absolutely everyone
- C) Top-down malicious use, where elites concentrate power and totalitarian backsliding occurs
- D) A lack of any single point of failure anywhere in the system as a whole

ANS: C) Top-down malicious use, where elites concentrate power and totalitarian backsliding occurs

8. What is an AI singleton?

- A) An open-source model that many people use freely
- B) An AI with decisive and lasting power over the world
- C) An AI system that has been permanently shut down
- D) A group of many AIs that are all equally powerful

ANS: B) An AI with decisive and lasting power over the world

9. What is a major cost of an AI singleton?

- A) It would increase the severity of collective action problems among competing developers
- B) It would increase competitive pressures among the developers of the AI systems
- C) It would be a single point of failure, so a rogue AI or malicious controller goes unopposed
- D) It would create many independent safeguards against failure across the whole system

ANS: C) It would be a single point of failure, so a rogue AI or malicious controller goes unopposed

10. Which is a cost of a diverse AI ecosystem?

- A) No AI would be able to interact with any other AI under any circumstances
- B) Power would always be distributed perfectly equally among all of the AIs involved
- C) Collective action problems would entirely disappear among the AIs in the ecosystem
- D) Multiple AIs could collude to emulate the failure modes of a singleton

ANS: D) Multiple AIs could collude to emulate the failure modes of a singleton

11. How does stakeholder theory differ from shareholder theory?

- A) It says only governments, and not shareholders, should be allowed to own corporations
- B) It says corporations should maximize value for their shareholders and for no one else
- C) It says corporations should ignore profit entirely in each and every decision they make
- D) It says corporations should also weigh employees, customers, partners, and society

ANS: D) It says corporations should also weigh employees, customers, partners, and society

12. Which legal form for a company can pursue public benefits rather than only maximizing shareholder value?

- A) Public benefit corporation (PBC)
- B) Shell company
- C) C corporation
- D) Sole proprietorship that issues no shares

ANS: A) Public benefit corporation (PBC)

13. In the ownership structure described in the deck, who elects the board of directors?

- A) Customers
- B) Shareholders
- C) Senior management
- D) Regulators

ANS: B) Shareholders

14. Why is legal liability described as a limited tool?

- A) A company could not cover all damages from a massive catastrophe, such as an AI-enabled bioweapon
- B) Liability has no effect on how companies behave in any circumstances at all
- C) Liability applies only to governments and never to companies, however large they are
- D) Liability can work only if strict government regulations are also already in place

ANS: A) A company could not cover all damages from a massive catastrophe, such as an AI-enabled bioweapon

15. What "tax loophole" does the deck point out regarding automation?

- A) Automated labor is taxed twice as heavily as human labor in most tax systems
- B) Only companies, and not individuals, pay any income tax on the labor they use
- C) AI systems are exempt from every kind of government regulation in every country
- D) Human labor is taxed (for example, income tax), but automated labor is not

ANS: D) Human labor is taxed (for example, income tax), but automated labor is not

16. Which of the following is listed as a way for a country to avoid falling behind while prioritizing safety?

- A) Publishing model weights openly so that rivals and allies alike can use them
- B) Avoiding all investment in information security for the country's own model weights
- C) Using immigration policy strategically to maintain a talent advantage
- D) Ending all forms of international cooperation on AI in order to move faster alone

ANS: C) Using immigration policy strategically to maintain a talent advantage

17. Which stage of international governance includes the IAEA conducting inspections of nuclear facilities?

- A) Policymaking by international bodies
- B) Agenda setting among the international community
- C) Evaluation, monitoring, and adjudication
- D) Implementation and enforcement of agreements

ANS: C) Evaluation, monitoring, and adjudication

18. Why does the deck say AI seems likely to be offense dominant?

- A) AI cannot be used to carry out any attacks on anything in the physical world
- B) It is easier to engineer pandemics with AI than to counteract them with AI
- C) Good AIs can always neutralize the threat that is posed by bad AIs in time
- D) Defensive systems are far cheaper and easier to build than offensive ones are

ANS: B) It is easier to engineer pandemics with AI than to counteract them with AI

19. Which property of compute makes it more governable than data or algorithms?

- A) It is intangible, so it can be copied at essentially no cost
- B) It is free to produce, so anyone can obtain more
- C) It cannot be quantified, so thresholds are impossible
- D) It is physical, so it is rivalrous and trackable

ANS: D) It is physical, so it is rivalrous and trackable

20. Which of the following is a compute security mechanism listed in the deck?

- A) Open publication of all of a model's training data to anyone who asks for it
- B) Unlimited remote access to the chips for any customer who is willing to pay
- C) Removing the chips' built-in ability to be tracked after they have been sold
- D) Location verification, which with offline licensing enables geofencing

ANS: D) Location verification, which with offline licensing enables geofencing